

BAX 452 HW 3

Course Instructor : Dr. Rahul Makhijani

Submit the assignment on canvas. The deadline is Tuesday Feb 11 11:59 pm

Bagging and Boosting on California Housing DataSet 20 points

1. Using the `RandomForestRegressor` from `sklearn.ensemble`, train a Random Forest model on the California Housing Dataset. Experiment with `n_estimators=100`, `random_state=42`. Train the model and print its default parameters.
2. **Tune the hyperparameters:**
 - **Task:** Use Grid Search or Random Search (`GridSearchCV` or `RandomizedSearchCV`) to optimize key hyperparameters such as:
 - Number of trees (`n_estimators`).
 - Maximum depth of trees (`max_depth`).
 - Minimum samples per leaf (`min_samples_leaf`).
 - **Question:** How does each parameter (above) influence the model's performance?
3. How does Gradient Boosting perform compared to other methods?
 - Create a residual plot (actual values vs. predicted values) to analyze the model's errors.
 - Are there any observable patterns in the prediction errors? Are errors larger for certain ranges of house prices?
 - Calculate the percentage error for high-price and low-price houses. Does the model perform equally well across the entire target range?
 - Observe how a very low or high value affects model accuracy, convergence speed, and overfitting.
4. **Experiment with the learning rate:**

- Observe how a very low or high value affects model accuracy, convergence speed, and overfitting.

IRIS multi-classifcation

Dataset Exploration and Preparation

- Select the *Iris* dafrom scikit-learn for classification. Load the dataset using `sklearn.datasets.load_*` functions.
- Display the dataset's feature names, target names, and a sample from the dataset. **Question:** Is the dataset well-balanced across the class labels? Comment on the distribution of target labels.

Data Preprocessing

- Split the dataset into training and testing datasets (80/20 split) using `train_test_split`.
- Normalize the feature values using `StandardScaler`.
- Explain the importance of feature scaling in KNN.

Implementing K-Nearest Neighbors

- Use scikit-learn's `KNeighborsClassifier` to train the KNN model.
- Train the model using the default parameters (`n_neighbors=5`, `metric='minkowski'`, `p=2` for Euclidean distance).
 - Fit the model on the training dataset and test the performance in terms of F1 score, precision and recall on the test set.
 - **Evaluate the impact of different values of k (n_neighbors):**
 - Train the model for different values of **k** ranging from 1 to 20.
 - Create a line plot of **k** vs. accuracy (for classification).
 - Find the optimal value of **k**.
 - **Question:** What value of **k** gives the best performance? Explain why the choice of **k** affects the model's performance.

Model Evaluation

- Generate a confusion matrix and a classification report (using `classification_report`) for `y_test` predictions.
- Interpret the confusion matrix and explain the precision, recall, and F1-score for each target class.
- Tune the KNN model's hyperparameters.
 - * Perform grid search or random search (`GridSearchCV`)
 - Number of neighbors (`n_neighbors`).
 - Distance metrics (`metric`, e.g., Euclidean, Manhattan).
 - Weighting schemes (`weights`, e.g., `uniform`, `distance`).
 - * **Question:** How does choosing different distance metrics (e.g., Euclidean vs. Manhattan) affect model performance?

Comparison with Other Algorithms

- * Compare KNN's performance with other models like:
 - Multinomial Logistic Regression.
 - Random Forest
- * Explore the impact of dropping certain features on the KNN model's performance.
- **Visualization:**
 - * For 2D datasets (e.g., *Iris*), visualize decision boundaries using `matplotlib` for different values of `k`.
 - * Comment on how decision boundaries change as `k` increases.

Implementing K-Means Clustering from Scratch (Iris Dataset)

The goal of this problem is to implement the K-Means clustering algorithm from scratch without relying on pre-built libraries like scikit-learn (except for loading or preprocessing the dataset). You will cluster data points into k groups using only the numeric features from the **Iris dataset**.

1. Dataset Preparation:

- * Load the Iris dataset from `sklearn.datasets`.
- * Use only the numeric features (*sepal length*, *sepal width*, *petal length*, *petal width*).
- * Normalize the features using a scaler (e.g., `StandardScaler` or `MinMaxScaler`).

2. K-Means Implementation:

- * Randomly initialize k centroids, choosing k data points randomly from the dataset. (**Note** this is slightly different from the initialization we had in class but this works as well)
- * Assign each data point to the nearest cluster centroid by calculating the Euclidean distance.
- * Update the cluster centroids to the mean of all points assigned to each cluster.
- * Repeat these steps until:
 - The cluster assignments do not change.
 - A maximum number of iterations (e.g., 100) is reached.
- * Implement the algorithm in Python, structuring it with reusable helper functions:
 - A function to calculate distances between points and centroids.
 - A function to assign clusters.
 - A function to update centroids.

3. Evaluation:

- * Use the **Elbow Method**:

- Apply your K-Means implementation for k ranging from 1 to 10.
- Calculate the total within-cluster variance for each value of k .
- Plot the total within-cluster variance values to identify the optimal number of clusters (elbow point).

4. **Cluster Visualization:**

- * Reduce the numeric data to 2D or 3D using **PCA** (Principal Component Analysis).
- * Visualize the clusters and centroids using a scatter plot.

5. **Analysis and Implementation Questions:**

- * What are the final centroids of the clusters for the chosen value of k ?
- * How does the within-cluster variance change as k increases?
- * Does your algorithm perform well for the Iris dataset? Why or why not?

Decision Tree

Consider a two-category classification task with the following training data:

attr1	attr2	attr3	attr4	class
a	1	c	-1	c1
b	0	c	-1	c1
a	0	c	1	c1
b	1	c	1	c1
b	0	c	1	c2
a	0	a	-1	c2
a	1	a	-1	c2
b	1	c	-1	c2

Table 1: Training Data for Two-Category Classification Task

1. Using the training data provided in Table 1, construct a complete, unpruned decision tree.
2. Use **information gain** as your splitting criterion for selecting the best attribute at each node.
3. Show all calculations for:
 - * The entropy of the dataset at each node.
 - * The weighted entropy for splits on each attribute.
 - * The information gain for each attribute.
4. Continue splitting until all examples in a node belong to the same class.

Hints and Definitions

- * The formula for **entropy** for a dataset D is:

$$H(D) = - \sum_{i=1}^C p_i \log_2(p_i)$$

where p_i is the proportion of examples belonging to class i , and C is the number of classes.

- * The **information gain** for a split on attribute A is:

$$IG(D, A) = H(D) - \sum_{v \in \text{values}(A)} \frac{|D_v|}{|D|} H(D_v)$$

where:

- $H(D)$ is the entropy of dataset D before the split.
 - D_v is the subset of D where attribute A has value v .
 - $|D_v|/|D|$ is the weight of subset D_v relative to the entire dataset.
 - $H(D_v)$ is the entropy of subset D_v .
- * Use the provided formulas for entropy and information gain to decide the best attribute to split at each step.